

# mnist-bharat-number-recognition

November 8, 2023

## 1 Recognising Hand written numbers

### Importing Libraries

```
[32]: import tensorflow
      from tensorflow import keras
      from tensorflow.keras import Sequential
      from tensorflow.keras.layers import Dense, Flatten
      import matplotlib.pyplot as plt
      from sklearn.metrics import accuracy_score
```

### downloading the dataset by dividing into train&test

```
[33]: (x_train,y_train), (x_test,y_test)= keras.datasets.mnist.load_data()
```

```
[34]: x_train.shape
```

```
[34]: (60000, 28, 28)
```

```
[35]: x_train
```

```
[35]: array([[0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          ...,
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0]],

          [[0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          ...,
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0]],

          [[0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          ...,
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0],
          [0, 0, 0, ..., 0, 0, 0]]]
```

```

[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]],

...,

[[0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 ...,
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 ...,
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0]], dtype=uint8)

```

```
[36]: x_test
```

```

[36]: array([[0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 ...,
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 ...,
 [0, 0, 0, ..., 0, 0, 0],

```

```

[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 ...,
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0]],

...,

[[0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 ...,
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 ...,
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0],
 [0, 0, 0, ..., 0, 0, 0]], dtype=uint8)


```

```
[37]: y_train
```

```
[37]: array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)
```


```
[38]: plt.imshow(x_train[1])
```

```
[38]: <matplotlib.image.AxesImage at 0x245295819a0>
```




```
[39]: plt.imshow(x_train[10])
```

```
[39]: <matplotlib.image.AxesImage at 0x2452ad62d30>
```



```
[70]: plt.imshow(x_train[140])
```

```
[70]: <matplotlib.image.AxesImage at 0x24517923730>
```



```
[41]: x_train[1] #It shows pixel values of first image. Generally its ranges from 0-255 so we need to normalize to 0-1 and train model
```

```
[41]: array([[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0],
               [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0],
               [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0],
               [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0],
               [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0, 51, 159, 253, 159, 50,  0,  0,  0,  0,  0,  0,
                0,  0],
               [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0, 48, 238, 252, 252, 252, 237,  0,  0,  0,  0,  0,  0,
                0,  0],
               [ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
                0,  0]
```

54, 227, 253, 252, 239, 233, 252, 57, 6, 0, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 10, 60,  
 224, 252, 253, 252, 202, 84, 252, 253, 122, 0, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 163, 252,  
 252, 252, 253, 252, 252, 96, 189, 253, 167, 0, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 51, 238, 253,  
 253, 190, 114, 253, 228, 47, 79, 255, 168, 0, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 48, 238, 252, 252,  
 179, 12, 75, 121, 21, 0, 0, 253, 243, 50, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 0, 0, 38, 165, 253, 233, 208,  
 84, 0, 0, 0, 0, 0, 0, 253, 252, 165, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 0, 7, 178, 252, 240, 71, 19,  
 28, 0, 0, 0, 0, 0, 0, 253, 252, 195, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 0, 57, 252, 252, 63, 0, 0,  
 0, 0, 0, 0, 0, 0, 0, 253, 252, 195, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 0, 198, 253, 190, 0, 0, 0,  
 0, 0, 0, 0, 0, 0, 0, 255, 253, 196, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 76, 246, 252, 112, 0, 0, 0,  
 0, 0, 0, 0, 0, 0, 0, 253, 252, 148, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 85, 252, 230, 25, 0, 0, 0,  
 0, 0, 0, 0, 0, 7, 135, 253, 186, 12, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 85, 252, 223, 0, 0, 0, 0,  
 0, 0, 0, 0, 7, 131, 252, 225, 71, 0, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 85, 252, 145, 0, 0, 0, 0,  
 0, 0, 0, 48, 165, 252, 173, 0, 0, 0, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 86, 253, 225, 0, 0, 0, 0,  
 0, 0, 114, 238, 253, 162, 0, 0, 0, 0, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 85, 252, 249, 146, 48, 29, 85,  
 178, 225, 253, 223, 167, 56, 0, 0, 0, 0, 0, 0, 0,  
 0, 0],  
 [ 0, 0, 0, 0, 0, 0, 85, 252, 252, 252, 229, 215, 252,  
 252, 252, 196, 130, 0, 0, 0, 0, 0, 0, 0, 0,  
 0, 0],

```
[ 0,  0,  0,  0,  0,  0, 28, 199, 252, 252, 253, 252, 252,
 233, 145,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0],
[ 0,  0,  0,  0,  0,  0,  0, 25, 128, 252, 253, 252, 141,
 37,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0],
[ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0,
  0,  0]], dtype=uint8)
```

### Normalising the pixel values to 0-1

```
[42]: X_train= x_train/255
      X_test= x_test/255
```

### Building a simple neural network

```
[43]: model=Sequential()
      model.add(Flatten(input_shape=(28,28)))
      model.add(Dense(128,activation='relu'))
      model.add(Dense(10,activation='softmax'))
      model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 784)	0
dense_2 (Dense)	(None, 128)	100480
dense_3 (Dense)	(None, 10)	1290

=====  
 Total params: 101,770  
 Trainable params: 101,770  
 Non-trainable params: 0  
 =====



```
[44]: model.  
      ↪ compile(loss='sparse_categorical_crossentropy', optimizer='Adam', metrics=['accuracy'])  
      model.fit(X_train, y_train, epochs=10, validation_split=0.2)
```

```
Epoch 1/10  
1500/1500 [=====] - 10s 6ms/step - loss: 0.2927 -  
accuracy: 0.9176 - val_loss: 0.1595 - val_accuracy: 0.9543  
Epoch 2/10  
1500/1500 [=====] - 9s 6ms/step - loss: 0.1327 -  
accuracy: 0.9614 - val_loss: 0.1294 - val_accuracy: 0.9634  
Epoch 3/10  
1500/1500 [=====] - 8s 5ms/step - loss: 0.0899 -  
accuracy: 0.9736 - val_loss: 0.1099 - val_accuracy: 0.9669  
Epoch 4/10  
1500/1500 [=====] - 7s 5ms/step - loss: 0.0661 -  
accuracy: 0.9803 - val_loss: 0.0966 - val_accuracy: 0.9706  
Epoch 5/10  
1500/1500 [=====] - 9s 6ms/step - loss: 0.0512 -  
accuracy: 0.9845 - val_loss: 0.0927 - val_accuracy: 0.9727  
Epoch 6/10  
1500/1500 [=====] - 10s 7ms/step - loss: 0.0397 -  
accuracy: 0.9881 - val_loss: 0.0888 - val_accuracy: 0.9743  
Epoch 7/10  
1500/1500 [=====] - 10s 7ms/step - loss: 0.0321 -  
accuracy: 0.9902 - val_loss: 0.1006 - val_accuracy: 0.9719  
Epoch 8/10  
1500/1500 [=====] - 9s 6ms/step - loss: 0.0250 -  
accuracy: 0.9924 - val_loss: 0.0882 - val_accuracy: 0.9756  
Epoch 9/10  
1500/1500 [=====] - 9s 6ms/step - loss: 0.0208 -  
accuracy: 0.9940 - val_loss: 0.0930 - val_accuracy: 0.9744  
Epoch 10/10  
1500/1500 [=====] - 8s 5ms/step - loss: 0.0169 -  
accuracy: 0.9953 - val_loss: 0.0957 - val_accuracy: 0.9747
```

```
[44]: <keras.callbacks.History at 0x2452ade18b0>
```

```
[45]: model.predict(X_test)
```

```
313/313 [=====] - 1s 3ms/step
```

```
[45]: array([[8.5839508e-10, 1.8917118e-12, 7.5510505e-09, ..., 9.9999893e-01,  
          2.4837241e-10, 2.8400832e-08],  
          [1.0145028e-09, 1.0243253e-06, 9.9999893e-01, ..., 4.0688998e-15,  
          9.2062621e-11, 2.5614173e-14],  
          [3.4333041e-04, 9.9575078e-01, 1.8154395e-04, ..., 2.6668913e-03,  
          6.5478130e-04, 5.6096642e-06],
```

```
...,
[7.9992448e-17, 1.2983661e-12, 1.4181497e-14, ..., 3.6875809e-07,
 2.3453754e-09, 1.0009860e-06],
[2.4981875e-15, 1.5609898e-15, 1.0015844e-15, ..., 2.4089571e-12,
 4.8702386e-08, 4.5870964e-14],
[5.7340820e-11, 2.1233560e-16, 4.0711426e-10, ..., 3.9116966e-16,
 4.8753010e-13, 3.9655894e-16]], dtype=float32)
```

```
[48]: y_prob=model.predict(X_test)
```

```
313/313 [=====] - 1s 3ms/step
```

```
[52]: y_prob
```

```
[52]: array([[8.5839508e-10, 1.8917118e-12, 7.5510505e-09, ..., 9.999893e-01,
 2.4837241e-10, 2.8400832e-08],
[1.0145028e-09, 1.0243253e-06, 9.999893e-01, ..., 4.0688998e-15,
 9.2062621e-11, 2.5614173e-14],
[3.4333041e-04, 9.9575078e-01, 1.8154395e-04, ..., 2.6668913e-03,
 6.5478130e-04, 5.6096642e-06],
...,
[7.9992448e-17, 1.2983661e-12, 1.4181497e-14, ..., 3.6875809e-07,
 2.3453754e-09, 1.0009860e-06],
[2.4981875e-15, 1.5609898e-15, 1.0015844e-15, ..., 2.4089571e-12,
 4.8702386e-08, 4.5870964e-14],
[5.7340820e-11, 2.1233560e-16, 4.0711426e-10, ..., 3.9116966e-16,
 4.8753010e-13, 3.9655894e-16]], dtype=float32)
```

```
[53]: y_prob.argmax(axis=1)
```

```
[53]: array([7, 2, 1, ..., 4, 5, 6], dtype=int64)
```

```
[54]: y_pred = y_prob.argmax(axis=1)
```

```
[55]: accuracy_score(y_test,y_pred)
```

```
[55]: 0.9756
```

```
[56]: history = model.fit(X_train,y_train,epochs=10,validation_split=0.2)
```

```
Epoch 1/10
```

```
1500/1500 [=====] - 9s 6ms/step - loss: 0.0138 -
accuracy: 0.9960 - val_loss: 0.0904 - val_accuracy: 0.9768
```

```
Epoch 2/10
```

```
1500/1500 [=====] - 9s 6ms/step - loss: 0.0112 -
accuracy: 0.9966 - val_loss: 0.0936 - val_accuracy: 0.9767
```

```
Epoch 3/10
```

```
1500/1500 [=====] - 9s 6ms/step - loss: 0.0114 -
```

```

accuracy: 0.9964 - val_loss: 0.0945 - val_accuracy: 0.9779
Epoch 4/10
1500/1500 [=====] - 9s 6ms/step - loss: 0.0079 -
accuracy: 0.9980 - val_loss: 0.0942 - val_accuracy: 0.9793
Epoch 5/10
1500/1500 [=====] - 9s 6ms/step - loss: 0.0091 -
accuracy: 0.9974 - val_loss: 0.1125 - val_accuracy: 0.9743
Epoch 6/10
1500/1500 [=====] - 9s 6ms/step - loss: 0.0062 -
accuracy: 0.9983 - val_loss: 0.1046 - val_accuracy: 0.9769
Epoch 7/10
1500/1500 [=====] - 9s 6ms/step - loss: 0.0065 -
accuracy: 0.9980 - val_loss: 0.1281 - val_accuracy: 0.9738
Epoch 8/10
1500/1500 [=====] - 9s 6ms/step - loss: 0.0067 -
accuracy: 0.9980 - val_loss: 0.1024 - val_accuracy: 0.9792
Epoch 9/10
1500/1500 [=====] - 10s 7ms/step - loss: 0.0048 -
accuracy: 0.9986 - val_loss: 0.1163 - val_accuracy: 0.9786
Epoch 10/10
1500/1500 [=====] - 9s 6ms/step - loss: 0.0055 -
accuracy: 0.9981 - val_loss: 0.1307 - val_accuracy: 0.9749

```

```
[57]: y_prob = model.predict(X_test)
```

```
313/313 [=====] - 1s 3ms/step
```


```
[58]: y_pred = y_prob.argmax(axis=1)
```

```
[59]: accuracy_score(y_test,y_pred)
```

```
[59]: 0.9757
```


```
[60]: plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
```

```
[60]: [<matplotlib.lines.Line2D at 0x24519603760>]
```



```
[61]: plt.plot(history.history['accuracy'])  
      plt.plot(history.history['val_accuracy'])
```

```
[61]: [<matplotlib.lines.Line2D at 0x24519bf1fa0>]
```



```
[62]: X_test
```

```
[62]: array([[0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            ...,
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.]],

          [[0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            ...,
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.]],

          [[0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            ...,
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.]])
```

```

[0., 0., 0., ..., 0., 0., 0.],
[0., 0., 0., ..., 0., 0., 0.],
[0., 0., 0., ..., 0., 0., 0.]],

...,

[[0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 ...,
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.]],

[[0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 ...,
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.]],

[[0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 ...,
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.],
 [0., 0., 0., ..., 0., 0., 0.]])


```

```
[63]: plt.imshow
```

```
[63]: <function matplotlib.pyplot.imshow(X, cmap=None, norm=None, *, aspect=None,
interpolation=None, alpha=None, vmin=None, vmax=None, origin=None, extent=None,
interpolation_stage=None, filternorm=True, filterrad=4.0, resample=None,
url=None, data=None, **kwargs)>
```

```
[64]: plt.imshow(X_test[0])
```

```
[64]: <matplotlib.image.AxesImage at 0x24517839df0>
```



```
[65]: model.predict(X_test[0].reshape(1,28,28))
```

```
1/1 [=====] - 0s 58ms/step
```

```
[65]: array([[3.7149667e-12, 2.4038387e-16, 1.2113532e-11, 6.8904434e-08,
            1.6952081e-22, 6.3673563e-12, 9.1037041e-18, 9.9999988e-01,
            9.5047105e-14, 1.2422003e-08]], dtype=float32)
```


```
[67]: model.predict(X_test[0].reshape(1,28,28)).argmax(axis=1)
```

```
1/1 [=====] - 0s 41ms/step
```

```
[67]: array([7], dtype=int64)
```

```
[68]: plt.imshow(X_test[1])
```


```
[68]: <matplotlib.image.AxesImage at 0x245178f12b0>
```



```
[71]: plt.imshow(X_test[6])
```

```
[71]: <matplotlib.image.AxesImage at 0x2451798c1c0>
```





```
[72]: model.predict(X_test[1].reshape(1,28,28))
```

```
1/1 [=====] - 0s 50ms/step
```

```
[72]: array([[9.0836498e-19, 2.2006896e-11, 1.0000000e+00, 1.1269086e-19,
            8.1333986e-28, 2.6898412e-24, 4.9036062e-20, 3.7776017e-24,
            9.2106045e-16, 8.8455555e-25]], dtype=float32)
```

```
[ ]:
```

```
[ ]:
```